



LLM performance: generalization benefit and latency trade-off

A. Accuracy by test family



B. Semantic decision cost



C. End-to-end control cost

